

CAS Advanced Machine Learning 2022

asdf

American Sign Language Recognition with DeepLearning

PhD Thesis submitted by

Felix Schlatter & Asad Bin Imtiaz

for the degree of

PhD in Cell Biology

Supervisor

Prof. Dr. Charles Darwin

asdf

asfd of the asdf

Co-advisor

asdf

Institute of Gregor Mendel

Faculty of Science of the University of Mendel

Table of contents

Table of contents	i
List of figures	iii
List of tables	v
1 Report	1
1 Introduction	1
2 Scope of the project	1
3 Dataset	1
4 Conclusion	2
5 References	2
A Data Augmentations	5
B Training Callbacks	7
1 Callbacks description	7
C Project Configuration	9
1 Project configuration description	9
D Data Utilities	15
1 Data processing Utils description	15
E HyperParameter Search	21
F Camera Stream Predictions	23
G ASL Dataset	25
1 ASL Dataset description	25
H Data Utilities	29
1 Deep Learning Utils	29
I Model Training	33
1 Generic trainer description	33
J Data Visualizations	37
K Pytorch Models	39
1 CVTransferLearningModel	41
1.1 Attributes	42
1.2 Methods	42
2 HybridEnsembleModel	42
2.1 Attributes	42

	2.2	Methods	43
3		HybridModel	43
	3.1	Attributes	43
	3.2	Methods	43
4		LSTMClassifier	44
	4.1	Attributes	44
	4.2	Methods	44
5		LSTMPredictor	44
	5.1	Attributes	44
	5.2	Methods	45
6		MultiHeadSelfAttention	45
	6.1	Attributes	45
	6.2	Methods	45
7		TransformerBlock	46
	7.1	Attributes	46
	7.2	Methods	46
8		TransformerEnsemble	46
	8.1	Attributes	46
	8.2	Methods	47
9		TransformerPredictor	47
	9.1	Attributes	47
	9.2	Methods	47
10		TransformerSequenceClassifier	48
	10.1	Attributes	48
	10.2	Methods	48
11		YetAnotherEnsemble	48
12		YetAnotherTransformer	49
	12.1	Attributes	49
	12.2	Methods	49
13		YetAnotherTransformerClassifier	50
	13.1	Attributes	50
	13.2	Methods	50

L	Tensorflow Models	51
----------	--------------------------	-----------

M	Indices and tables	57
----------	---------------------------	-----------

Bibliography	59
---------------------	-----------

Bibliography	59
---------------------	-----------

Python Module Index	61
----------------------------	-----------

Python Module Index	61
----------------------------	-----------

Index	62
--------------	-----------

List of figures

1	This is the caption for the image.	2
2	Figure caption goes here.	3

List of tables

1	Model Classes	39
---	-------------------------	----

Chapter 1

Report

1 Introduction

This is the introduction section of the report. Provide an overview of the report's purpose, scope, and key objectives.

2 Scope of the project

The objectives of this project were dual pronged. The first aim was to leverage the power of Deep Learning to extract valuable insights from an existing dataset. Concurrently, we strived to avoid binding our project to a single Deep Learning framework. Our approach was focused on establishing a versatile boilerplate setup conducive for the development of models using both PyTorch and TensorFlow frameworks. This strategy enabled us to delve into and exploit the unique capabilities and subtleties of each platform, ensuring our work was framework-agnostic and compatible with the most widely used Deep Learning technologies.

By implementing the boilerplate setup, we aimed to streamline the process of deploying Deep Learning models in real-world scenarios, regardless of the framework preference of the end-users. The goal was to provide a unified and easily adaptable codebase that could be utilized by others working with Deep Learning models, simplifying the transition from development to deployment. Through this project's scope, we sought to advance our understanding of Deep Learning concepts, maximize the potential of the dataset, and contribute to the broader Deep Learning community by providing a practical solution for model deployment using PyTorch and TensorFlow.

In the spirit of creating a user-friendly and dynamic setup, we centralized the control over various inputs through the `config.py` file. This configuration file offers a one-stop solution for defining and tweaking project-specific parameters, settings, and variables. Users can readily customize the project's behavior and characteristics by simply adjusting the values in the `config.py` file, thereby eliminating the need to directly modify the code.

The project involved rigorous data preprocessing, cleaning, and augmentation to tailor the dataset for the specific use case. The setup was designed to be scalable, capable of handling larger datasets or more complex models. It was built with a view to accommodate future developments in deep learning and sign language recognition. One key aspect of the scope was to create reusable modules and functions to ensure that parts of the project can be easily used in other similar tasks, thereby reducing future development time. Thereby we entailed detailed documentation to ensure that other researchers and developers can easily understand and extend the work done.

3 Dataset

The dataset used for this project was provided by Google and hosted by Kaggle and was part of the 2023 challenge Isolated Sign Language Recognition [1][3]. This rich and comprehensive dataset has been a joint contribution from Google, the Georgia Institute of Technology, the National Technical Institute for the Deaf at Rochester Institute of Technology, and Deaf Professional Arts Network. The data comprised a total of 54.43 GB of processed data. split into 94'479 unique files containing one of 250 distinct signs. Each sign is performed by different participants, thus capturing a wide variety of signing styles, personal quirks, and potential minor variations in sign performance. This arrangement allows for the capture of a broad spectrum of signing styles, personal idiosyncrasies, and possible minor discrepancies in sign execution. The breadth and depth of this dataset make it an invaluable resource for understanding and modeling sign language patterns. The videos are preprocessed using the MediaPipe -Holistic pipeline [4] and are thus represented as a series of 543 distinct landmark coordinates (x,y,z). The landmarks comprise coordinates for pose, both hands, and face. This representation not only helps reduce the data complexity but also enables efficient and accurate analysis of the sign language patterns.



Fig. 1: This is the caption for the image.

This is the content of section 2. .. raw:: latex
blindtext blindtext blindtext

4 Conclusion

Summarize the main findings and conclusions of the report.

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for

special content, but the length of words should match the language. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

[1].

5 References

clearpage section*{References}
clearpage section{References}
clearpage section{References}
clearpage section{References}

+ + G + P
INGENIEURE

Fig. 2: Figure caption goes here.

Appendix A

Data Augmentations

`augmentations.frame_dropout(frames, dropout_rate=0.05)`

Randomly drop frames from the input landmark data.

Args:

`frames` (numpy.ndarray): An array of landmarks data. `dropout_rate` (float): The proportion of frames to drop (default: 0.05).

Returns:

numpy.ndarray: An array of landmarks with dropped frames.

`augmentations.mirror_landmarks(frames)`

Invert/mirror landmark coordinates along the x-axis.

Args:

`frames` (numpy.ndarray): An array of landmarks data.

Returns:

numpy.ndarray: An array of inverted landmarks.

`augmentations.normalize(frames, mn, std)`

Normalize the frames with a given mean and standard deviation.

Args:

`frames` (numpy.ndarray): An array of landmarks data. `mn` (float): The mean value for normalization. `std` (float): The standard deviation for normalization.

Returns:

numpy.ndarray: An array of normalized landmarks.

`augmentations.random_rotation(frames, max_angle=10)`

Apply random rotation to landmark coordinates. (on X and Y only)

Args:

`frames` (numpy.ndarray): An array of landmarks data. `max_angle` (int): The maximum rotation angle in degrees (default: 10).

Returns:

numpy.ndarray: An array of landmarks with randomly rotated coordinates.

`augmentations.random_scaling(frames, scale_range=(0.9, 1.1))`

Apply random scaling to landmark coordinates.

Args:

`frames` (numpy.ndarray): An array of landmarks data. `scale_range` (tuple): A tuple containing the minimum and maximum scaling factors (default: (0.9, 1.1)).

Returns:

numpy.ndarray: An array of landmarks with randomly scaled coordinates.

`augmentations.shift_landmarks(frames, max_shift=0.01)`

Shift landmark coordinates randomly by a small amount.

Args:

`frames` (numpy.ndarray): An array of landmarks data. `max_shift` (float): Maximum shift for the random shift (default: 0.01).

Returns:

numpy.ndarray: An array of augmented landmarks.

`augmentations.standardize(frames)`

Standardize the frames so that they have mean 0 and standard deviation 1.

Args:

`frames` (numpy.ndarray): An array of landmarks data.

Returns:

numpy.ndarray: An array of standardized landmarks.

Appendix B

Training Callbacks

1 Callbacks description

This module contains callback codes which may be executed during training. These callbacks are used to dynamically adjust the dropout rate and data augmentation probability during the training process, which can be useful techniques to prevent overfitting and increase the diversity of the training data, potentially improving the model's performance.

The `dropout_callback` function is designed to increase the dropout rate of the model during the training process after a certain number of epochs. The dropout rate is a regularization technique used to prevent overfitting during the training process. The rate of dropout is increased every few epochs based on a specified rate until it reaches a specified maximum limit.

The `augmentation_increase_callback` function is designed to increase the probability of data augmentation applied to the dataset during the training process after a certain number of epochs. Data augmentation is a technique that can generate new training samples by applying transformations to the existing data. The probability of data augmentation is increased every few epochs based on a specified rate until it reaches a specified maximum limit.

```
callbacks.augmentation_increase_callback(trainer, aug_increase_rate=1.5, max_limit=0.35)
```

A callback function designed to increase the probability of data augmentation applied on the dataset during the training process. Data augmentation is a technique that can generate new training samples by applying transformations to the existing data.

The increase in data augmentation is performed every few epochs based on 'DYNAMIC_AUG_INC_INTERVAL' until it reaches a specified maximum limit.

Args:

`trainer`: The object that contains the model and handles the training process. `aug_increase_rate`: The rate at which data augmentation probability is increased. Default is value of 'DYNAMIC_AUG_INC_RATE' from config. `max_limit`: The maximum limit to which data augmentation probability can be increased. Default is value of 'DYNAMIC_AUG_MAX_THRESHOLD' from config.

Returns:

None

Functionality:

Increases the probability of data augmentation applied on the dataset after certain number of epochs defined by 'DYNAMIC_AUG_INC_INTERVAL'.

Parameters

- `trainer` – Trainer object handling the training process
- `aug_increase_rate` – Rate at which to increase the data augmentation probability
- `max_limit` – Maximum allowable data augmentation probability

Returns

None

Return type

None

```
callbacks.dropout_callback(trainer, dropout_rate=1.1, max_dropout=0.2)
```

A callback function designed to increase the dropout rate of the model in training after a certain number of epochs. The dropout rate is a regularization technique which helps in preventing overfitting during the training process.

The rate of dropout is increased every few epochs based on the config parameter (in config.py) 'DYNAMIC_DROP_OUT_REDUCTION_INTERVAL' until a maximum threshold defined by 'max_dropout'. This function is usually called after each epoch in the training process.

Args:

trainer: The object that contains the model and handles the training process.
dropout_rate: The rate at which the dropout rate is increased. Default is value of 'DYNAMIC_DROP_OUT_REDUCTION_RATE' from config.
max_dropout: The maximum limit to which dropout can be increased. Default is value of 'DYNAMIC_DROP_OUT_MAX_THRESHOLD' from config.

Returns:

None

Functionality:

Increases the dropout rate of all nn.Dropout modules in the model after certain number of epochs defined by 'DYNAMIC_DROP_OUT_REDUCTION_INTERVAL'.

Parameters

- *trainer* – Trainer object handling the training process
- *dropout_rate* – Rate at which to increase the dropout rate
- *max_dropout* – Maximum allowable dropout rate

Returns

None

Return type

None

Appendix C

Project Configuration

1 Project configuration description

This configuration is created allows for easy tuning of your machine learning model's parameters and setup. The device on which the model runs, the paths for various resources, the seed for random number generation, hyperparameters for model training, and much more and quickly be change and configured. This makes your setup flexible and easy to adapt for various experiments and environments

```
config.BATCH_SIZE = 128
```

Training Batch Size

```
config.CHECKPOINT_DIR = 'checkpoints/'
```

Checkpoint files Directory path

```
config.CLEANED_FILE = 'cleansed_data.marker'
```

File that marks the data cleaning stage.

```
config.COLUMNS_TO_USE = ['x', 'y']
```

Coordinate columns from the data to use for training.

```
config.DATA_DIR = 'data/'
```

Data files Directory path

```
config.DEVICE = 'cpu'
```

Setting the device for training, 'cuda' if a CUDA-compatible GPU is available, 'mps' if multiple processors are available, 'cpu' if none of the above.

```
config.DL_FRAMEWORK = 'tensorflow'
```

Deep learning framework to use for training and inference. Can be either 'pytorch' or 'tensorflow'.

```
config.DYNAMIC_AUG_INC_INTERVAL = 5
```

The number of epochs to wait before increasing the probability of data augmentation.

```
config.DYNAMIC_AUG_INC_RATE = 1.5
```

The rate at which the probability of data augmentation is increased.

```
config.DYNAMIC_AUG_MAX_THRESHOLD = 0.35
```

The maximum limit to which the probability of data augmentation can be increased.

```
config.DYNAMIC_DROP_OUT_INIT_RATE = 0.01
```

The value of initial low dropouts rate

```
config.DYNAMIC_DROP_OUT_MAX_THRESHOLD = 0.2
```

The max value of dynamic dropouts

`config.DYNAMIC_DROP_OUT_REDUCTION_INTERVAL = 2`

The epoch interval value to gradually change dropout rate

`config.DYNAMIC_DROP_OUT_REDUCTION_RATE = 1.1`

The value to increase dropouts by

`config.EARLY_STOP_METRIC = 'accuracy'`

Which metric should be used for early stopping loss/accuracy

`config.EARLY_STOP_MODE = 'max'`

What is the mode? min/max

`config.EARLY_STOP_PATIENCE = 10`

The number of epochs to wait for improvement in the validation loss before stopping training

`config.EARLY_STOP_TOLERANCE = 0.001`

The value of loss as margin to tolerate

`config.EPOCHS = 50`

Training Number of epochs

`config.FACE_FEATURES = 468`

Number of features related to the face in the data.

`config.FACE_FEATURE_START = 0`

Start index for face feature in the data.

`config.FACE_INDICES = array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39], dtype=int64)`

Indices of face landmarks that are used from the data.

`config.FACE_LANDMARKS = array([61, 185, 40, 39, 37, 0, 267, 269, 270, 409, 291, 146, 91, 181, 84, 17, 314, 405, 321, 375, 78, 191, 80, 81, 82, 13, 312, 311, 310, 415, 95, 88, 178, 87, 14, 317, 402, 318, 324, 308])`

Landmarks for Lips

`config.HAND_FEATURES = 21`

Number of features related to the hand in the data.

`config.HAND_INDICES = array([40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81], dtype=int64)`

Indices of hand landmarks that are used from the data.

`config.INPUT_SIZE = 32`

Size of the input data for the model.

`config.INTERMOLATE_MISSING = 3`

Number of missing values to interpolate in the data.

`config.LANDMARK_FILES = 'train_landmark_files'`

Directory where training landmark files are stored.

`config.LEARNING_RATE = 0.001`

Training Learning rate

`config.LEFT_HAND_FEATURE_START = 468`

Start index for left hand feature in the data.


```
config.LEFT_HAND_INDICES = array([40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54,
55, 56, 57, 58, 59, 60], dtype=int64)
```

Indices of left hand landmarks that are used from the data.

```
config.MAP_JSON_FILE = 'sign_to_prediction_index_map.json'
```

JSON file that maps sign to prediction index.

```
config.MARKER_FILE = 'preprocessed_data.marker'
```

File that marks the preprocessing stage.

```
config.MAX_SEQUENCES = 32
```

Maximum number of sequences in the input data.

```
config.MIN_SEQUEENCES = 8.0
```

Minimum number of sequences in the input data.

```
config.MODELNAME = 'CVTransferLearningModel'
```

Name of the model to be used for training.

```
config.MODEL_DIR = 'models/'
```

Model files Directory path

```
config.N_CLASSES = 250
```

Number of classes

```
config.N_DIMS = 2
```

Number of dimensions used in training

```
config.N_LANDMARKS = 96
```

Total number of used landmarks

```
config.OUT_DIR = 'out/'
```

Output files Directory path

```
config.POSE_FEATURES = 33
```

Number of features related to the pose in the data.

```
config.POSE_FEATURE_START = 489
```

Start index for pose feature in the data.

```
config.POSE_INDICES = array([82, 83, 84, 85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95],
dtype=int64)
```

Indices of pose landmarks that are used from the data.

```
config.PROCESSED_DATA_DIR = 'data/processed/'
```

Processed Data files Directory path

```
config.RAW_DATA_DIR = 'data/raw/'
```

Raw Data files Directory path

```
config.RIGHT_HAND_FEATURE_START = 522
```

Start index for right hand feature in the data.

```
config.RIGHT_HAND_INDICES = array([61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75,
76, 77, 78, 79, 80, 81], dtype=int64)
```

Indices of right hand landmarks that are used from the data.

```
config.ROOT_PATH = 'C:\\Users\\fs.GUNDP\\Python\\CAS-AML-FINAL-PROJECT\\docs\\../src\\../'
```

Root directory

```
config.ROWS_PER_FRAME = 543
```

Number of rows per frame in the data.

```
config.RUNS_DIR = 'runs/'
    Run files Directory path

config.SEED = 0
    Set Random Seed

config.SKIP_CONSECUTIVE_ZEROS = 4
    Skip data if there are this many consecutive zeros.

config.SRC_DIR = 'src/'
    Source files Directory path

config.TEST_SIZE = 0.05
    Testing Test set size

config.TRAIN_CSV_FILE = 'train.csv'
    CSV file name that contains the training dataset.

config.TRAIN_SIZE = 0.9
    Training Train set split size

config.TUNE_HP = True
    Tune hyperparameters

config.USED_FACE_FEATURES = 40
    Count of facial features used

config.USED_HAND_FEATURES = 21
    Count of hands features used (single hand only)

config.USED_POSE_FEATURES = 14
    Count of body/pose features used

config.USEFUL_ALL_LANDMARKS = array([ 61, 185, 40, 39, 37, 0, 267, 269, 270, 409, 291, 146,
91, 181, 84, 17, 314, 405, 321, 375, 78, 191, 80, 81, 82, 13, 312, 311, 310, 415, 95, 88, 178,
87, 14, 317, 402, 318, 324, 308, 468, 469, 470, 471, 472, 473, 474, 475, 476, 477, 478, 479,
480, 481, 482, 483, 484, 485, 486, 487, 488, 522, 523, 524, 525, 526, 527, 528, 529, 530, 531,
532, 533, 534, 535, 536, 537, 538, 539, 540, 541, 542, 500, 501, 502, 503, 504, 505, 506, 507,
508, 509, 510, 511, 512, 513])
    All Landmarks

config.USEFUL_FACE_LANDMARKS = array([ 61, 185, 40, 39, 37, 0, 267, 269, 270, 409, 291, 146,
91, 181, 84, 17, 314, 405, 321, 375, 78, 191, 80, 81, 82, 13, 312, 311, 310, 415, 95, 88, 178,
87, 14, 317, 402, 318, 324, 308])
    Landmarks for face

config.USEFUL_HAND_LANDMARKS = array([468, 469, 470, 471, 472, 473, 474, 475, 476, 477, 478,
479, 480, 481, 482, 483, 484, 485, 486, 487, 488, 522, 523, 524, 525, 526, 527, 528, 529, 530,
531, 532, 533, 534, 535, 536, 537, 538, 539, 540, 541, 542])
    Landmarks for both hands

config.USEFUL_LEFT_HAND_LANDMARKS = array([468, 469, 470, 471, 472, 473, 474, 475, 476, 477,
478, 479, 480, 481, 482, 483, 484, 485, 486, 487, 488])
    Landmarks for left hand

config.USEFUL_POSE_LANDMARKS = array([500, 501, 502, 503, 504, 505, 506, 507, 508, 509, 510,
511, 512, 513])
    Landmarks for pose

config.USEFUL_RIGHT_HAND_LANDMARKS = array([522, 523, 524, 525, 526, 527, 528, 529, 530, 531,
532, 533, 534, 535, 536, 537, 538, 539, 540, 541, 542])
    Landmarks for right hand
```

```
config.VALID_SIZE = 0.05
```

Training Validation set size

Appendix D

Data Utilities

1 Data processing Utils description

This module handles the loading and preprocessing of data. It is specifically tailored for loading ASL sign language dataset where the raw data includes information about the position of hands, face, and body over time.

ASL stands for American Sign Language, which is a natural language used by individuals who are deaf or hard of hearing to communicate through hand gestures and facial expressions.

The dataset consists of sequences of frames, where each frame contains multiple “landmarks”. Each of these landmarks has multiple features, such as coordinates. The landmarks may represent various aspects of human body, such as facial features, hand positions, and body pose.

This module is used to process the raw data, to create a uniform dataset where all sequences are of the same length and all missing values have been handled in a way that maintains the integrity of the data. This involves steps like detecting and removing empty frames, selecting specific landmarks, resizing sequences and handling NaN values.

`data.data_utils.calculate_avg_landmark_positions(dataset)`

Calculate the average landmark positions for left-hand, right-hand, and face landmarks for each sign in the dataset. The purpose of this function is to compute the average positions of landmarks for left-hand, right-hand, and face for each sign in the training dataset.

Returns: List : Containing a dictionary with average x/y positions with keys - ‘left_hand’ - ‘right_hand’ - ‘face’

Functionality: - The function takes an ASLDataset object as an input, which contains the training data. - It calculates the average landmark positions for left-hand, right-hand, and face landmarks for each sign in the dataset. - The function returns a list containing a dictionary with average x/y positions with keys ‘left_hand’, ‘right_hand’, and ‘face’ for each sign.

Parameters

`dataset` ([ASL_DATASET](#)) – The ASL dataset object containing the training data.

Returns

A list containing a dictionary with average x/y positions with keys ‘left_hand’, ‘right_hand’, and ‘face’

for each sign. :rtype: List[Dict[str, np.ndarray]]

`data.data_utils.calculate_landmark_length_stats()`

Calculate statistics of landmark lengths for each sign type.

Returns: dict: A dictionary of landmark lengths for each sign type containing: - minimum - maximum - mean - median - standard deviation

Functionality: - The function reads the CSV file. - It groups the DataFrame by sign. - An empty dictionary is created to store average landmarks for each sign type. - The function loops through each unique sign and its corresponding rows in the grouped DataFrame. - For each sign, it initializes a list to store the length of

landmarks for each example of the current sign. - It loops through each row of the current sign type, loads the data, and adds the length of landmarks of the current example to the list of current sign data. - The function calculates the minimum, maximum, mean, standard deviation, and median of the landmarks for the current sign and updates the dictionary. - The resulting dictionary containing average landmarks for each sign type is returned.

Returns

A dictionary of landmark lengths for each sign type containing minimum, maximum, mean, median & standard

deviation :rtype: dict

```
data.data_utils.create_data_loaders(asl_dataset, train_size=0.9, valid_size=0.05, test_size=0.05,
                                    batch_size=128, random_state=0, dl_framework='tensorflow',
                                    num_workers=16)
```

Split the ASL dataset into training, validation, and testing sets and create data loaders for each set.

Args: *asl_dataset* (*ASLDataset*): The ASL dataset to load data from. *train_size* (float, optional): The proportion of the dataset to include in the training set. Defaults to 0.8. *valid_size* (float, optional): The proportion of the dataset to include in the validation set. Defaults to 0.1. *test_size* (float, optional): The proportion of the dataset to include in the testing set. Defaults to 0.1. *batch_size* (int, optional): The number of samples per batch to load. Defaults to BATCH_SIZE. *random_state* (int, optional): The seed used by the random number generator for shuffling the data. Defaults to SEED.

Returns: tuple of *Loader*: A tuple containing the data loaders for training, validation, and testing sets.

Parameters

- *asl_dataset* (*ASLDataset*) – The ASL dataset to load data from.
- *train_size* (*float*) – The proportion of the dataset to include in the training set.
- *valid_size* (*float*) – The proportion of the dataset to include in the validation set.
- *test_size* (*float*) – The proportion of the dataset to include in the testing set.
- *batch_size* (*int*) – The number of samples per batch to load.
- *random_state* (*int*) – The seed used by the random number generator for shuffling the data.

Returns

A tuple containing the data loaders for training, validation, and testing sets.

Return type

tuple of *Loader*

```
data.data_utils.interpolate_missing_values(arr, max_gap=3)
```

This function provides a solution for handling missing values in the data array. It interpolates these missing values, filling them with plausible values that maintain the overall data integrity. The function uses a linear interpolation method that assumes a straight line between the two points on either side of the gap. The maximum gap size for which interpolation should be performed is also configurable.

Arguments: The function takes two arguments - an array with missing values, and a maximum gap size for interpolation. If the size of the gap (i.e., number of consecutive missing values) is less than or equal to this specified maximum gap size, the function will fill it with interpolated values. This ensures that the data maintains its continuity without making too far-fetched estimations for larger gaps.

Args:

arr (*np.ndarray*): Input array with missing values. *max_gap* (*int*, optional): Maximum gap to fill. Defaults to INTEREMOLATE_MISSING.

Returns:

np.ndarray: Array with missing values interpolated.

Functionality:

Interpolates missing values in the array. The function fills gaps of up to a maximum size with interpolated values, maintaining data integrity and continuity.

Returns

Array with missing values interpolated.

Return type

np.ndarray

Parameters

- `arr` (*np.ndarray*) – Input array with missing values.
- `max_gap` (*int*) – Maximum gap to fill.

This function uses linear interpolation to fill the missing values. Other forms of interpolation such as polynomial or spline may provide better results for specific types of data. It is also worth noting that no imputation method can fully recover original data, and as such, results should be interpreted with caution when working with imputed data.

data.data_utils.load_relevant_data_subset(pq_path)

This function serves a key role in handling data in our pipeline by loading only a subset of the relevant data from a given path. The primary purpose of this is to reduce memory overhead when working with large datasets. The implementation relies on efficient data loading strategies, leveraging the speed of Parquet file format and the ability to read in only necessary chunks of data instead of the whole dataset.

The function takes as input a string which represents the path to the data file. It makes use of pandas' parquet read function to read the data file. This function is particularly suited for reading large datasets as it allows for efficient on-disk storage and fast query capabilities. The function uses PyArrow library as the engine for reading the parquet files which ensures efficient and fast reading of data. After reading the data, the function selects the relevant subset based on certain criteria, which is task specific.

Args:

`pq_path` (str): Path to the data file.

Returns:

np.ndarray: Subset of the relevant data as a NumPy array.

Functionality:

Loads a subset of the relevant data from a given path.

Returns

Subset of the relevant data.

Return type

np.ndarray

Parameters

`pq_path` (*str*) – Path to the data file.

The function assumes that the data file is in parquet format and the necessary libraries for reading parquet files are installed. It also assumes that the path provided is a valid path to the data file.

data.data_utils.preprocess_data(landmarks)

This function preprocesses the input data by applying similar steps as the `preprocess_data_to_same_size` function, but with the difference that it does not interpolate missing values. The function again targets to adjust the size of the input data to align with the `INPUT_SIZE`. It selects only non-empty frames and follows similar strategies of padding, repeating, and pooling the data for size alignment.

Args:

`landmarks` (np.ndarray): The input array with landmarks data.

Returns:

Tuple[np.ndarray, int]: A tuple containing processed landmark data and the final size of the data.

Parameters

`landmarks` (*np.ndarray*) – The input array with landmarks data.

Returns

A tuple containing processed landmark data and the final size of the data.

Return type

Tuple[np.ndarray, int]

`data.data_utils.preprocess_data_item(raw_landmark_path, targets_sign)`

The function preprocesses landmark data for a single file. The process involves applying transformations to raw landmark data to convert it into a form more suitable for machine learning models. The transformations may include normalization, scaling, etc. The target sign associated with the landmark data is also taken as input.

This function is a handy function to process all landmark sequences on a particular location. This will come in handy while testing where individual sequences may be provided

Args:

`raw_landmark_path`: Path to the raw landmark file
`targets_sign`: The target sign for the given landmark data

Returns: dict: A dictionary containing the preprocessed landmarks, target, and size.

Functionality: - The function reads the parquet file and processes the data. - It filters columns to include only frame, type, landmark_index, x, and y. - The function then filters face mesh landmarks and pose landmarks based on the predefined useful landmarks. - Landmarks data is pivoted to have a multi-level column structure on landmark type and frame sequence ids. - Missing values are interpolated using linear interpolation, and any remaining missing values are filled with 0. - The function rearranges columns and calculates the number of frames in the data. - X and Y coordinates are brought together, and a dictionary with the processed data is created and returned.

Parameters

- `raw_landmark_path` (*str*) – Path to the raw landmark file.
- `targets_sign` (*int*) – The target sign for the given landmark data.

Returns

A dictionary containing the preprocessed landmarks, target, and size.

Return type

dict

`data.data_utils.preprocess_data_to_same_size(landmarks)`

This function preprocesses the input data to ensure all data arrays have the same size, specified by the global `INPUT_SIZE` variable. This uniform size is necessary for subsequent processing and analysis stages, particularly those involving machine learning models which often require consistent input sizes. The preprocessing involves several steps, including handling missing values, upsampling, and reshaping arrays. It begins by interpolating any missing values, and then it subsets the data by selecting only non-empty frames. Various strategies are applied to align the data size to the desired `INPUT_SIZE`, including padding, repeating, and pooling the data.

Args:

`landmarks` (`np.ndarray`): The input array with landmarks data.

Returns:

Tuple[`np.ndarray`, int, int, int]: A tuple containing processed landmark data, the set input size, the number of original frames, and the number of frames after preprocessing.

Parameters

`landmarks` (`np.ndarray`) – The input array with landmarks data.

Returns

A tuple containing processed landmark data, the set input size, the number of original frames, and the

number of frames after preprocessing. :rtype: Tuple[`np.ndarray`, int, int, int]

`data.data_utils.preprocess_raw_data(sample=100000)`

Preprocesses the raw data, saves it as numpy arrays into processed data directory and updates the metadata CSV file.

This method `preprocess_data` preprocesses the data for easier and faster loading during training time. The data is processed and stored in `PROCESSED_DATA_DIR` if not already done.

This function is responsible for preprocessing raw data. The primary functionality involves converting raw data into a format more suitable for the machine learning pipeline, namely NumPy arrays. The function operates

on a sample of data, allowing for efficient processing of large datasets in manageable chunks. Additionally, this function also takes care of persisting the preprocessed data for future use and updates the metadata accordingly.

Args: sample (int): Number of samples to preprocess. Default is 100000.

Functionality: - The function reads the metadata CSV file for training data to obtain a dictionary that maps target values to integer indices. - It then reads the training data CSV file and generates the absolute path to locate landmark files. - Next, it keeps text signs and their respective indices and initializes a list to store the processed data. - The data is then processed and stored in the list by iterating over each file path in the training data and reading in the parquet file for that file path. - The landmark data is then processed and padded to have a length of max_seq_length. - Finally, a dictionary with the processed data is created and added to the list. - The processed data is saved to disk using the np.save method and the saved file is printed.

Parameters

sample (*int*, *optional*, *default: 100000*) – Number of samples to preprocess.

Returns

None

Note: If the preprocessed data already exists, the function prints “Preprocessed data found. Skipping...” and exits.

`data.data_utils.remove_outlier_or_missing_data(landmark_len_dict)`

This function removes rows from the training data that contain missing or outlier landmark data. It takes as input a dictionary containing the statistics of landmark lengths for each sign type. The function processes the training data and removes rows with missing or outlier landmark data. The function also includes a nested function ‘has_consecutive_zeros’ which checks for consecutive frames where X and Y coordinates are both zero. If a cleansing marker file exists, it skips the process, indicating that the data is already cleaned.

Functionality:

This function takes a dictionary with the statistics of landmark lengths per sign type and uses it to identify outlier sequences. It removes any rows with missing or outlier landmark data. An outlier sequence is defined as one that is either less than a third of the median length or more than two standard deviations away from the mean length. A row is also marked for deletion if the corresponding landmark file is missing or if the sign’s left-hand or right-hand landmarks contain more than a specified number of consecutive zeros.

Args:

landmark_len_dict (dict): A dictionary containing the statistics of landmark lengths for each sign type.

Returns:

None

Parameters

landmark_len_dict (*dict*) – A dictionary containing the statistics of landmark lengths for each sign type.

Returns

None, the function doesn’t return anything. It modifies data in-place.

`data.data_utils.remove_unusable_data()`

This function checks the existing training data for unusable instances, like missing files or data that is smaller than the set minimum sequence length. If unusable data is found, it is removed from the system, both in terms of files and entries in the training dataframe. The dataframe is updated and saved back to the disk. If a cleansing marker file exists, it skips the process, indicating that the data is already cleaned.

Functionality:

The function iterates through the DataFrame rows, attempting to load and check each landmark file specified in the row’s path. If the file is missing or if the file’s usable size is less than a predefined threshold, the function deletes the corresponding landmark file and marks the row for deletion in the DataFrame. At the end, the function removes all marked rows from the DataFrame, updates it and saves it to the disk.

Returns:

None

Returns

None, the function doesn't return anything. It modifies data in-place.

Appendix E

HyperParameter Search

Appendix F

Camera Stream Predictions

```
predict_on_camera.convert_mp_to_df(results)  
predict_on_camera.show_camera_feed(model, LAST_FRAMES=32)
```


Appendix G

ASL Dataset

1 ASL Dataset description

This file contains the `ASL_DATASET` class which serves as the dataset module for American Sign Language (ASL) data. The `ASL_DATASET` is designed to load, preprocess, augment, and serve the dataset for model training and validation. This class provides functionalities such as loading the dataset from disk, applying transformations, data augmentation techniques, and an interface to access individual data samples.

Note: This dataset class expects data in a specific format. Detailed explanations and expectations about input data are provided in respective method docstrings.

```
class data.dataset.ASL_DATASET(metadata_df=None, transform=None, max_seq_length=32,
                                augment=False, augmentation_threshold=0.1, enableDropout=True)
```

A dataset class for the ASL dataset.

The `ASL_DATASET` class represents a dataset of American Sign Language (ASL) gestures, where each gesture corresponds to a word or phrase. This class provides functionalities to load the dataset, apply transformations, augment the data, and yield individual data samples for model training and validation.

`__getitem__(idx)`

Get an item from the dataset by index.

This method returns a data sample from the dataset based on a provided index. It handles reading of the processed data file, applies transformations and augmentations (if set), and pads the data to match the maximum sequence length. It returns the preprocessed landmarks and corresponding target as a tuple.

Args:

idx (int): The index of the item to retrieve.

Returns:

tuple: A tuple containing the landmarks and target for the item.

Functionality:

Get a single item from the dataset.

Parameters

idx (*int*) – The index of the item to retrieve.

Returns

A tuple containing the landmarks and target for the item.

Return type

tuple

```
__init__(metadata_df=None, transform=None, max_seq_length=32, augment=False,
          augmentation_threshold=0.1, enableDropout=True)
```

Initialize the ASL dataset.

This method initializes the dataset and loads the metadata necessary for the dataset processing. If no metadata is provided, it will load the default processed dataset. It also sets the transformation functions, data augmentation parameters, and maximum sequence length.

Args:

`metadata_df` (`pd.DataFrame`, optional): A dataframe containing the metadata for the dataset. Defaults to None. `transform` (callable, optional): A function/transform to apply to the data. Defaults to None. `max_seq_length` (`int`, optional): The maximum sequence length for the data. Defaults to `INPUT_SIZE`. `augment` (`bool`, optional): Whether to apply data augmentation. Defaults to False. `augmentation_threshold` (`float`, optional): Probability of augmentation happening. Only if `augment == True`. Defaults to 0.1. `enableDropout` (`bool`, optional): Whether to enable the frame dropout augmentation. Defaults to True.

Functionality:

Initializes the dataset with necessary configurations and loads the data.

Parameters

- `metadata_df` (*`pd.DataFrame`*, *optional*) – A dataframe containing the metadata for the dataset.
- `transform` (*callable*, *optional*) – A function/transform to apply to the data.
- `max_seq_length` (*int*) – The maximum sequence length for the data.
- `augment` (*bool*) – Whether to apply data augmentation.
- `augmentation_threshold` (*float*) – Probability of augmentation happening. Only if `augment == True`.
- `enableDropout` (*bool*) – Whether to enable the frame dropout augmentation.

`__len__()`

Get the length of the dataset.

This method returns the total number of data samples present in the dataset. It's an implementation of the special method `__len__` in Python, providing a way to use the Python built-in function `len()` on the dataset object.

Functionality:

Get the length of the dataset.

Returns:

`int`: The length of the dataset.

Returns

The length of the dataset.

Return type

`int`

`__repr__()`

Return a string representation of the ASL dataset.

This method returns a string that provides an overview of the dataset, including the number of participants and total data samples. It's an implementation of the special method `__repr__` in Python, providing a human-readable representation of the dataset object.

Returns:

`str`: A string representation of the dataset.

Functionality:

Return a string representation of the dataset.

Returns

A string representation of the dataset.

Return type

`str`

`--weakref--`

list of weak references to the object (if defined)

`load_data()`

Load the data for the ASL dataset.

This method loads the actual ASL data based on the metadata provided during initialization. If no metadata was provided, it loads the default processed data. It generates absolute paths to locate landmark files, and stores individual metadata lists for easy access during data retrieval.

Functionality:

Loads the data for the dataset.

Return type

None

Appendix H

Data Utilities

1 Deep Learning Utils

This module provides a set of helper functions that abstract away specific details of different deep learning frameworks (such as TensorFlow and PyTorch). These functions allow the main code to run in a framework-agnostic manner, thus improving code portability and flexibility.

```
class dl_utils.DatasetWithLen(tf_dataset, length)
```

The DatasetWithLen class serves as a wrapper around TensorFlow's Dataset object. Its primary purpose is to add a length method to the TensorFlow Dataset. This is useful in contexts where it's necessary to know the number of batches that a DataLoader will create from a dataset, which is a common requirement in many machine learning training loops. It also provides an iterator over the dataset, which facilitates traversing the dataset for operations such as batch creation.

For instance, this might be used in conjunction with a progress bar during training to display the total number of batches. Since TensorFlow's Dataset objects don't inherently have a `__len__` method, this wrapper class provides that functionality, augmenting the dataset with additional features that facilitate the training process.

Args:

`tf_dataset`: The TensorFlow dataset to be wrapped. `length`: The length of the dataset.

Functionality:

Provides a length method and an iterator for a TensorFlow dataset.

Return type

DatasetWithLen object

Parameters

- `tf_dataset` – The TensorFlow dataset to be wrapped.
- `length` – The length of the dataset.

```
__init__(tf_dataset, length)
```

```
__iter__()
```

Returns an iterator for the dataset.

Returns

iterator for the dataset

```
__len__()
```

Returns the length of the dataset.

Returns

length of the dataset

`--weakref--`

list of weak references to the object (if defined)

`dl_utils.get_PT_Dataset(dataloader)`

Retrieve the underlying dataset from a PyTorch data loader.

Parameters

`dataloader` – DataLoader object.

Returns

Dataset object.

`dl_utils.get_TF_Dataset(dataloader)`

Retrieve the underlying dataset from a TensorFlow data loader.

Parameters

`dataloader` – DatasetWithLen object.

Returns

Dataset object.

`dl_utils.get_dataloader(dataset, batch_size=128, shuffle=True, dl_framework='tensorflow',
num_workers=16)`

The `get_dataloader` function is responsible for creating a `DataLoader` object given a dataset and a few other parameters. A `DataLoader` is an essential component in machine learning projects as it controls how data is fed into the model during training. However, different deep learning frameworks have their own ways of creating and handling `DataLoader` objects.

To improve the portability and reusability of the code, this function abstracts away these specifics, allowing the user to create a `DataLoader` object without having to worry about the details of the underlying framework (TensorFlow or PyTorch). This approach can save development time and reduce the risk of bugs or errors.

Args:

`dataset`: The dataset to be loaded. `batch_size`: The size of the batches that the `DataLoader` should create. `shuffle`: Whether to shuffle the data before creating batches. `dl_framework`: The name of the deep learning framework. `num_workers`: The number of worker threads to use for loading data.

Functionality:

Creates and returns a `DataLoader` object that is compatible with the specified deep learning framework.

Return type

`DataLoader` or `DatasetWithLen` object

Parameters

- `dataset` – The dataset to be loaded.
- `batch_size` – The size of the batches that the `DataLoader` should create.
- `shuffle` – Whether to shuffle the data before creating batches.
- `dl_framework` – The name of the deep learning framework.
- `num_workers` – The number of worker threads to use for loading data.

`dl_utils.get_dataset(dataloader)`

The `get_dataset` function is an interface to extract the underlying dataset from a `dataloader`, irrespective of the deep learning framework being used, i.e., TensorFlow or PyTorch. The versatility of this function makes it integral to any pipeline designed to be flexible across both TensorFlow and PyTorch frameworks.

Given a `dataloader` object, this function first determines the deep learning framework currently in use by referring to the `DL_FRAMEWORK` config parameter variable. If the framework is TensorFlow, it invokes the `get_TF_Dataset` function to retrieve the dataset. Alternatively, if PyTorch is being used, the `get_PT_Dataset` function is called. This abstracts away the intricacies of handling different deep learning frameworks, thereby simplifying the process of working with datasets across TensorFlow and PyTorch.

Args:

`dataloader`: `DataLoader` from PyTorch or `DatasetWithLen` from TensorFlow.

Functionality:

Extracts the underlying dataset from a `dataloader`, be it from PyTorch or TensorFlow.

Return type

Dataset object

Parameters`dataloader` – DataLoader in case of PyTorch and DatasetWithLen in case of TensorFlow.`dl_utils.get_model_params(model_name)`

The `get_model_params` function is a utility function that serves to abstract away the details of reading model configurations from a YAML file. In a machine learning project, it is common to have numerous models, each with its own set of hyperparameters. These hyperparameters can be stored in a YAML file for easy access and modification.

This function reads the configuration file and retrieves the specific parameters associated with the given model. The configurations are stored in a dictionary which is then returned. This aids in maintaining a cleaner, more organized codebase and simplifies the process of updating or modifying model parameters.

Args:`model_name`: Name of the model whose parameters are to be retrieved.**Functionality:**

Reads a YAML file and retrieves the model parameters as a dictionary.

Return type

dict

Parameters`model_name` – Name of the model whose parameters are to be retrieved.`dl_utils.log_metrics(writer, log_dict)`

Helper function to log metrics to TensorBoard.

Parameters

- `log_dict` – Dictionary to log all the metrics to tensorboard. It must contain the keys {epoch, accuracy, loss, lr,}
- `writer` – TensorBoard writer object.

Type`log_dict``dl_utils.to_PT_DataLoader(dataset, batch_size=128, shuffle=True, num_workers=16)`

This function is the PyTorch counterpart to ‘`to_TF_DataLoader`’. It converts a given dataset into a PyTorch DataLoader. The purpose of this function is to streamline the creation of PyTorch DataLoaders, allowing for easy utilization in a PyTorch training or inference pipeline.

The PyTorch DataLoader handles the process of drawing batches of data from a dataset, which is essential when training models. This function further extends this functionality by implementing data shuffling and utilizing multiple worker threads for asynchronous data loading, thereby optimizing the data loading process during model training.

Args:

`dataset`: The dataset to be loaded. `batch_size`: The size of each batch the DataLoader will return. `shuffle`: Whether the data should be shuffled before batching. `num_workers`: The number of worker threads to use for data loading.

Functionality:

Converts a given dataset into a PyTorch DataLoader.

Return type

DataLoader object

Parameters

- `dataset` – The dataset to be loaded.
- `batch_size` – The size of each batch the DataLoader will return.
- `shuffle` – Whether the data should be shuffled before batching.
- `num_workers` – The number of worker threads to use for data loading.

`dl_utils.to_TF_DataLoader(dataset, batch_size=128, shuffle=True)`

This function takes in a dataset and converts it into a TensorFlow DataLoader. Its purpose is to provide a streamlined method to generate DataLoaders that can be utilized in a TensorFlow training or inference pipeline. It not only ensures the dataset is in a format that can be ingested by TensorFlow's pipeline, but also implements optional shuffling of data, which is a common practice in model training to ensure random distribution of data across batches.

This function first checks whether the data is already in a tensor format, if not it converts the data to a tensor. Next, it either shuffles the dataset or keeps it as is, based on the 'shuffle' flag. Lastly, it prepares the TensorFlow DataLoader by batching the dataset and applying an automatic optimization strategy for the number of parallel calls in mapping functions.

Args:

dataset: The dataset to be loaded. batch_size: The size of each batch the DataLoader will return.
shuffle: Whether the data should be shuffled before batching.

Functionality:

Converts a given dataset into a TensorFlow DataLoader.

Return type

DatasetWithLen object

Parameters

- **dataset** – The dataset to be loaded.
- **batch_size** – The size of each batch the DataLoader will return.
- **shuffle** – Whether the data should be shuffled before batching.

Appendix I

Model Training

1 Generic trainer description

Trainer module handles the training, validation, and testing of framework-agnostic deep learning models.

The Trainer class handles the complete lifecycle of model training including setup, execution of training epochs, validation and testing, early stopping, and result logging.

The class uses configurable parameters for defining training settings like early stopping and batch size, and it supports adding custom callback functions to be executed at the end of each epoch. This makes the trainer class flexible and adaptable for various types of deep learning models and tasks.

Attributes: `model_name` (str): The name of the model to be trained. `params` (dict): The parameters required for the model. `model` (model object): The model object built using the given model name and parameters. `train_loader`, `valid_loader`, `test_loader` (DataLoader objects): PyTorch dataloaders for training, validation, and testing datasets. `patience` (int): The number of epochs to wait before stopping training when the validation loss is no longer improving. `best_val_metric` (float): The best validation metric recorded. `patience_counter` (int): A counter that keeps track of the number of epochs since the validation loss last improved. `model_class` (str): The class name of the model. `train_start_time` (str): The starting time of the training process. `writer` (SummaryWriter object): TensorBoard's SummaryWriter to log metrics for visualization. `checkpoint_path` (str): The path where the best model checkpoints will be saved during training. `epoch` (int): The current epoch number. `callbacks` (list): A list of callback functions to be called at the end of each epoch.

Methods: `train(n_epochs)`: Trains the model for a specified number of epochs. `evaluate()`: Evaluates the model on the validation set. `test()`: Tests the model on the test set. `add_callback(callback)`: Adds a callback function to the list of functions to be called at the end of each epoch.

```
class trainer.Trainer(modelname='CVTransferLearningModel', dataset=<class
                        'data.dataset.ASL_DATASET'>, patience=10, enableAugmentationDropout=True,
                        augmentation_threshold=0.35)
```

A trainer class which acts as a control hub for the model lifecycle, including initial setup, executing training epochs, performing validation and testing, implementing early stopping, and logging results. The module has been designed to be agnostic to the specific deep learning framework, enhancing its versatility across various projects.

```
__init__(modelname='CVTransferLearningModel', dataset=<class 'data.dataset.ASL_DATASET'>,
          patience=10, enableAugmentationDropout=True, augmentation_threshold=0.35)
```

Initializes the Trainer class with the specified parameters.

This method initializes various components needed for the training process. This includes the model specified by the model name, the dataset with optional data augmentation and dropout, data loaders for the training, validation, and test sets, a SummaryWriter for logging, and a path for saving model checkpoints.

1. The method first retrieves the specified model and its parameters.

2. It then initializes the dataset and the data loaders.
3. It sets up metrics for early stopping and a writer for logging.
4. Finally, it prepares a directory for saving model checkpoints.

Args:

`modelname` (*str*): The name of the model to be used for training. `dataset` (*Dataset*): The dataset to be used. `patience` (*int*): The number of epochs with no improvement after which training will be stopped. `enableAugmentationDropout` (*bool*): If True, enable data augmentation dropout. `augmentation_threshold` (*float*): The threshold for data augmentation.

Functionality:

This method initializes various components, such as the model, dataset, data loaders, logging writer, and checkpoint path, required for the training process.

Parameters

- `modelname` (*str*) – The name of the model for training.
- `dataset` (*Dataset*) – The dataset for training.
- `patience` (*int*) – The number of epochs with no improvement after which training will be stopped.
- `enableAugmentationDropout` (*bool*) – If True, enable data augmentation dropout.
- `augmentation_threshold` (*float*) – The threshold for data augmentation.

Return type

None

Note: This method only initializes the Trainer class. The actual training is done by calling the `train()` method.

Warning: Make sure the specified model name corresponds to an actual model in your project's models directory.

`__weakref__`

list of weak references to the object (if defined)

`add_callback(callback)`

Adds a callback to the Trainer.

This method simply appends a callback function to the list of callbacks stored by the Trainer instance. These callbacks are called at the end of each training epoch.

Functionality:

It allows the addition of custom callbacks to the training process, enhancing its flexibility.

Parameters

- `callback` (*Callable*) – The callback function to be added.

Returns

None

Return type

None

Warning: The callback function must be callable and should not modify the training process.

`evaluate()`

Evaluates the model on the validation set.

This method sets the model to evaluation mode and loops over the validation dataset, computing the loss and accuracy for each batch. It then averages these metrics and logs them. This process provides an unbiased estimate of the model's performance on new data during training.

Functionality:

It manages the evaluation of the model on the validation set, handling batch-wise loss computation and accuracy assessment.

Returns

Average validation loss and accuracy

Return type

Tuple[float, float]

Warning: Ensure the model is in evaluation mode to correctly compute the validation metrics.

test()

Tests the model on the test set.

This method loads the best saved model, sets it to evaluation mode, and then loops over the test dataset, computing the loss, accuracy, and predictions for each batch. It then averages the loss and accuracy and logs them. It also collects all the model's predictions and their corresponding labels.

Functionality:

It manages the testing of the model on the test set, handling batch-wise loss computation, accuracy assessment, and prediction generation.

Returns

List of all predictions and their corresponding labels

Return type

Tuple[List, List]

train(*n_epochs*=50, *limit_batches*=None)

Trains the model for a specified number of epochs.

This method manages the main training loop of the model. For each epoch, it performs several steps. It first puts the model into training mode and loops over the training dataset, calculating the loss and accuracy for each batch and optimizing the model parameters. It logs these metrics and updates a progress bar. At the end of each epoch, it evaluates the model on the validation set and checks whether early stopping criteria have been met. If the early stopping metric has improved, it saves the current model and its parameters. If not, it increments a counter and potentially stops training if the counter exceeds the allowed patience. Finally, it steps the learning rate scheduler and calls any registered callbacks.

1. The method first puts the model into training mode and initializes some lists and counters.
2. Then it enters the main loop over the training data, updating the model and logging metrics.
3. It evaluates the model on the validation set and checks the early stopping criteria.
4. If the criteria are met, it saves the model and its parameters; if not, it increments a patience counter.
5. It steps the learning rate scheduler and calls any callbacks.

Args:

n_epochs (int): The number of epochs for which the model should be trained.

Functionality:

This method coordinates the training of the model over a series of epochs, handling batch-wise loss computation, backpropagation, optimization, validation, early stopping, and model checkpoint saving.

Parameters

n_epochs (*int*) – Number of epochs for training.

Returns

None

Return type

None

Note: This method modifies the state of the model and its optimizer, as well as various attributes of the Trainer instance itself.

Warning: If you set the patience value too low in the constructor, the model might stop training prematurely.

Appendix J

Data Visualizations

`visualizations.visualize_data_distribution(dataset)`

Visualize the distribution of data in terms of the number of samples and average sequence length per class.

This function generates two bar charts: one showing the number of samples per class, and the other showing the average sequence length per class.

Parameters

`dataset` (*ASL_Dataset*) – The ASL dataset to load data from.

`visualizations.visualize_target_sign(dataset, target_sign, n_samples=6)`

Visualize *n_samples* instances of a given target sign from the dataset.

This function generates a visual representation of the landmarks for each sample belonging to the specified *target_sign*.

Args:

`dataset` (*ASL_Dataset*): The ASL dataset to load data from. `target_sign` (*int*): The target sign to visualize. `n_samples` (*int*, optional): The number of samples to visualize. Defaults to 6.

Returns:

`matplotlib.animation.FuncAnimation`: A matplotlib animation object displaying the landmarks for each frame.

Parameters

- `dataset` (*ASL_Dataset*) – The ASL dataset to load data from.
- `target_sign` (*int*) – The target sign to visualize.
- `n_samples` (*int*, *optional*) – The number of samples to visualize, defaults to 6.

Returns

A matplotlib animation object displaying the landmarks for each frame.

Return type

`matplotlib.animation.FuncAnimation`

Appendix K

Pytorch Models

This module defines a PyTorch *BaseModel* providing a basic framework for learning and validating from *Trainer* module, from which other pytorch models are inherited. This module includes several model classes that build upon the PyTorch's nn.Module for constructing pytorch LSTM or Transformer based models:

Table 1: Model Classes

Class	Description
<i>TransformerSequenceClassifier</i>	This is a transformer-based sequence classification model. The class constructs a transformer encoder based on user-defined parameters or default settings. The forward method first checks and reshapes the input, then passes it through the transformer layers. It then pools the sequence by taking the mean over the time dimension, and finally applies the output layer to generate the class predictions.
<i>TransformerPredictor</i>	A <i>TransformerPredictor</i> model that extends the Pytorch <i>BaseModel</i> . This class wraps <i>TransformerSequenceClassifier</i> model and provides functionality to use it for making predictions.
<i>MultiHeadSelfAttention</i>	This class applies a multi-head attention mechanism. It has options for causal masking and layer normalization. The input is expected to have dimensions [batch_size, seq_len, features].
<i>TransformerBlock</i>	This class represents a single block of a transformer architecture, including multi-head self-attention and a feed-forward neural network, both with optional layer normalization and dropout. The input is expected to have dimensions [batch_size, seq_len, features].
<i>YetAnotherTransformerClassifier</i>	This class constructs a transformer-based classifier with a specified number of <i>TransformerBlock</i> instances. The output of the model is a tensor of logits with dimensions [batch_size, num_classes].
<i>YetAnotherTransformer</i>	This class is a wrapper for <i>YetAnotherTransformerClassifier</i> which includes learning rate, optimizer, and learning rate scheduler settings. It extends from the <i>BaseModel</i> class.
<i>YetAnotherEnsemble</i>	This class constructs an ensemble of <i>YetAnotherTransformerClassifier</i> instances, where the outputs are concatenated and passed through a fully connected layer. This class also extends from the <i>BaseModel</i> class and includes learning rate, optimizer, and learning rate scheduler settings.

```
class models.pytorch.models.BaseModel(*args: Any, **kwargs: Any)
```

A *BaseModel* that extends the nn.Module from PyTorch.

Functionality: #. The class initializes with a given learning rate and number of classes. #. It sets up the loss criterion, accuracy metric, and default states for optimizer and scheduler. #. It defines an abstract method 'forward' which should be implemented in the subclass. #. It also defines various utility functions like calculating accuracy, training, validation and testing steps, scheduler stepping, and model checkpointing.

Args:

`learning_rate` (float): The initial learning rate for optimizer. `n_classes` (int): The number of classes for classification.

Parameters

- `learning_rate` (*float*) – The initial learning rate for optimizer.
- `n_classes` (*int*) – The number of classes for classification.

Returns

None

Return type

None

Note: The class does not directly initialize the optimizer and scheduler. They should be initialized in the subclass if needed.

Warning: The ‘forward’ function must be implemented in the subclass, else it will raise a `NotImplementedError`.

`__init__(learning_rate, n_classes=250)`

`calculate_accuracy(y_hat, y)`

Calculates the accuracy of the model’s prediction.

Parameters

- `y_hat` (*Tensor*) – The predicted output from the model.
- `y` (*Tensor*) – The ground truth or actual labels.

Returns

The calculated accuracy.

Return type

Tensor

`eval_mode()`

Sets the model to evaluation mode.

`forward(x)`

The forward function for the BaseModel.

Parameters

- `x` (*Tensor*) – The inputs to the model.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a `NotImplementedError`.

`get_lr()`

Gets the current learning rate of the model.

Returns

The current learning rate.

Return type

float

`load_checkpoint(filepath)`

Loads the model and optimizer states from a checkpoint.

Parameters

- `filepath` (*str*) – The file path where to load the model checkpoint from.

`optimize()`

Steps the optimizer and sets the gradients of all optimized `torch.Tensor` s to zero.

`save_checkpoint(filepath)`

Saves the model and optimizer states to a checkpoint.

Parameters

`filepath (str)` – The file path where to save the model checkpoint.

`step_scheduler()`

Steps the learning rate scheduler, adjusting the optimizer's learning rate as necessary.

`test_step(batch)`

Performs a test step using the input batch data.

Parameters

`batch (tuple)` – A tuple containing input data and labels.

Returns

The calculated loss, accuracy, and model predictions.

Return type

tuple

`train_mode()`

Sets the model to training mode.

`training_step(batch)`

Performs a training step using the input batch data.

Parameters

`batch (tuple)` – A tuple containing input data and labels.

Returns

The calculated loss and accuracy.

Return type

tuple

`validation_step(batch)`

Performs a validation step using the input batch data.

Parameters

`batch (tuple)` – A tuple containing input data and labels.

Returns

The calculated loss and accuracy.

Return type

tuple

```
class models.pytorch.models.CVTransferLearningModel(*args: Any, **kwargs: Any)
```

1 CVTransferLearningModel

A CVTransferLearningModel that extends the Pytorch BaseModel.

This class applies transfer learning for computer vision tasks using pretrained models. It also provides a forward method to pass an input through the model.

1.1 Attributes

learning_rate

[float] The learning rate for the optimizer.

model

[nn.Module] The base model for transfer learning.

optimizer

[torch.optim.Adam] The optimizer used for updating the model parameters.

scheduler

[torch.optim.lr_scheduler.ExponentialLR] The learning rate scheduler used for adapting the learning rate during training.

1.2 Methods

forward(x)

Performs a forward pass through the model.

__init__(kwargs)****forward(x)**

The forward function for the BaseModel.

Parameters

x (*Tensor*) – The inputs to the model.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a NotImplementedError.

```
class models.pytorch.models.HybridEnsembleModel(*args: Any, **kwargs: Any)
```

2 HybridEnsembleModel

A HybridEnsembleModel that extends the Pytorch BaseModel.

This class creates an ensemble of LSTM and Transformer models and provides functionality to use the ensemble for making predictions.

2.1 Attributes

learning_rate

[float] The learning rate for the optimizer.

lstms

[nn.ModuleList] The list of LSTM models.

models

[nn.ModuleList] The list of Transformer models.

fc

[nn.Linear] The final fully-connected layer.

optimizer

[torch.optim.Adam] The optimizer used for updating the model parameters.

scheduler

[torch.optim.lr_scheduler.ExponentialLR] The learning rate scheduler used for adapting the learning rate during training.

2.2 Methods

forward(x)

Performs a forward pass through the model.

`__init__(**kwargs)`

forward(x)

The forward function for the BaseModel.

Parameters

`x` (*Tensor*) – The inputs to the model.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a `NotImplementedError`.

```
class models.pytorch.models.HybridModel(*args: Any, **kwargs: Any)
```

3 HybridModel

A HybridModel that extends the Pytorch BaseModel.

This class combines the LSTMClassifier and TransformerSequenceClassifier models and provides functionality to use the combined model for making predictions.

3.1 Attributes

lstm

[LSTMClassifier] The LSTM classifier used for making predictions.

transformer

[TransformerSequenceClassifier] The transformer sequence classifier used for making predictions.

fc

[nn.Linear] The final fully-connected layer.

optimizer

[torch.optim.Adam] The optimizer used for updating the model parameters.

scheduler

[torch.optim.lr_scheduler.ExponentialLR] The learning rate scheduler used for adapting the learning rate during training.

3.2 Methods

forward(x)

Performs a forward pass through the model.

`__init__(**kwargs)`

forward(x)

The forward function for the BaseModel.

Parameters

`x` (*Tensor*) – The inputs to the model.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a `NotImplementedError`.

```
class models.pytorch.models.LSTMClassifier(*args: Any, **kwargs: Any)
```

4 LSTMClassifier

A LSTM-based Sequence Classifier. This class utilizes a LSTM network for sequence classification tasks.

4.1 Attributes

DEFAULTS

[dict] Default settings for the LSTM and classifier. These can be overridden by passing values in the constructor.

lstm

[nn.LSTM] The LSTM network used for processing the input sequence.

dropout

[nn.Dropout] The dropout layer applied after LSTM network.

output_layer

[nn.Linear] The output layer used to generate class predictions.

4.2 Methods

forward(x)

Performs a forward pass through the model.

__init__(kwargs)**

forward(x)

Forward pass through the model

```
class models.pytorch.models.LSTMPredictor(*args: Any, **kwargs: Any)
```

5 LSTMPredictor

A LSTMPredictor model that extends the Pytorch BaseModel.

This class wraps the LSTMClassifier model and provides functionality to use it for making predictions.

5.1 Attributes

learning_rate

[float] The learning rate for the optimizer.

model

[LSTMClassifier] The LSTM classifier used for making predictions.

optimizer

[torch.optim.Adam] The optimizer used for updating the model parameters.

scheduler

[torch.optim.lr_scheduler.ExponentialLR] The learning rate scheduler used for adapting the learning rate during training.

5.2 Methods

forward(x)

Performs a forward pass through the model.

`__init__(**kwargs)`

forward(x)

The forward function for the BaseModel.

Parameters

`x` (*Tensor*) – The inputs to the model.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a `NotImplementedError`.

```
class models.pytorch.models.MultiHeadSelfAttention(*args: Any, **kwargs: Any)
```

6 MultiHeadSelfAttention

A MultiHeadSelfAttention module that extends the `nn.Module` from PyTorch.

Functionality: `#`. The class initializes with a given dimension size, number of attention heads, dropout rate, layer normalization and causality. `#`. It sets up the multihead attention module and layer normalization. `#`. It also defines a forward method that applies the multihead attention, causal masking if requested, and layer normalization if requested.

6.1 Attributes

multihead_attn

[`nn.MultiheadAttention`] The multihead attention module.

layer_norm

[`nn.LayerNorm` or `None`] The layer normalization module. If it is not applied, set to `None`.

causal

[`bool`] If `True`, applies causal masking.

6.2 Methods

forward(x)

Performs a forward pass through the model.

Args:

`dim` (`int`): The dimension size of the input data. `num_heads` (`int`): The number of attention heads. `dropout` (`float`): The dropout rate. `layer_norm` (`bool`): Whether to apply layer normalization. `causal` (`bool`): Whether to apply causal masking.

Returns: `None`

```
__init__(dim, num_heads=8, dropout=0.1, layer_norm=True, causal=True)
```

```
class models.pytorch.models.TransformerBlock(*args: Any, **kwargs: Any)
```

7 TransformerBlock

A TransformerBlock module that extends the nn.Module from PyTorch.

Functionality: `#`. The class initializes with a given dimension size, number of attention heads, expansion factor, attention dropout rate, and dropout rate. `#`. It sets up the multihead self-attention module, layer normalization and feed-forward network. `#`. It also defines a forward method that applies the multihead self-attention, dropout, layer normalization and feed-forward network.

7.1 Attributes

norm1, norm2, norm3

[nn.LayerNorm] The layer normalization modules.

attn

[MultiHeadSelfAttention] The multihead self-attention module.

feed_forward

[nn.Sequential] The feed-forward network.

dropout

[nn.Dropout] The dropout module.

7.2 Methods

forward(x)

Performs a forward pass through the model.

Args:

`dim` (int): The dimension size of the input data. `num_heads` (int): The number of attention heads. `expansion_factor` (int): The expansion factor for the hidden layer size in the feed-forward network. `attn_dropout` (float): The dropout rate for the attention module. `drop_rate` (float): The dropout rate for the module.

Returns: None

```
--init--(dim=192, num_heads=4, expansion_factor=4, attn_dropout=0.2, drop_rate=0.2)
```

```
class models.pytorch.models.TransformerEnsemble(*args: Any, **kwargs: Any)
```

8 TransformerEnsemble

A TransformerEnsemble that extends the Pytorch BaseModel.

This class creates an ensemble of TransformerSequenceClassifier models and provides functionality to use the ensemble for making predictions.

8.1 Attributes

learning_rate

[float] The learning rate for the optimizer.

models

[nn.ModuleList] The list of transformer sequence classifiers used for making predictions.

fc

[nn.Linear] The final fully-connected layer.

optimizer

[torch.optim.Adam] The optimizer used for updating the model parameters.

scheduler

[torch.optim.lr_scheduler.ExponentialLR] The learning rate scheduler used for adapting the learning rate during training.

8.2 Methods

forward(x)

Performs a forward pass through the model.

`__init__(**kwargs)`

forward(x)

The forward function for the BaseModel.

Parameters

`x` (*Tensor*) – The inputs to the model.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a `NotImplementedError`.

```
class models.pytorch.models.TransformerPredictor(*args: Any, **kwargs: Any)
```

9 TransformerPredictor

A TransformerPredictor model that extends the Pytorch BaseModel.

This class wraps the TransformerSequenceClassifier model and provides functionality to use it for making predictions.

9.1 Attributes

learning_rate

[float] The learning rate for the optimizer.

model

[TransformerSequenceClassifier] The transformer sequence classifier used for making predictions.

optimizer

[torch.optim.Adam] The optimizer used for updating the model parameters.

scheduler

[torch.optim.lr_scheduler.ExponentialLR] The learning rate scheduler used for adapting the learning rate during training.

9.2 Methods

forward(x)

Performs a forward pass through the model.

`__init__(**kwargs)`

forward(x)

The forward function for the BaseModel.

Parameters

`x` (*Tensor*) – The inputs to the model.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a `NotImplementedError`.

```
class models.pytorch.models.TransformerSequenceClassifier(*args: Any, **kwargs: Any)
```

10 TransformerSequenceClassifier

A Transformer-based Sequence Classifier. This class utilizes a transformer encoder to process the input sequence.

The transformer encoder consists of a stack of N transformer layers that are applied to the input sequence. The output sequence from the transformer encoder is then passed through a linear layer to generate class predictions.

10.1 Attributes

DEFAULTS

[dict] Default settings for the transformer encoder and classifier. These can be overridden by passing values in the constructor.

transformer

[nn.TransformerEncoder] The transformer encoder used to process the input sequence.

output_layer

[nn.Linear] The output layer used to generate class predictions.

batch_first

[bool] Whether the first dimension of the input tensor represents the batch size.

10.2 Methods

forward(inputs)

Performs a forward pass through the model.

```
__init__(**kwargs)
```

forward(inputs)

Forward pass through the model

```
class models.pytorch.models.YetAnotherEnsemble(*args: Any, **kwargs: Any)
```

11 YetAnotherEnsemble

A YetAnotherEnsemble model that extends the Pytorch BaseModel.

Functionality: #. The class initializes with a set of parameters for the YetAnotherTransformerClassifier. #. It sets up an ensemble of YetAnotherTransformerClassifier models, a fully connected layer, the optimizer and the learning rate scheduler. #. It also defines a forward method that applies each YetAnotherTransformerClassifier model in the ensemble, concatenates the outputs and applies the fully connected layer.

Args:

kwargs (dict): A dictionary containing the parameters for the YetAnotherTransformerClassifier models, fully connected layer, optimizer and learning rate scheduler.

Returns: None

```
__init__(**kwargs)
```

```
forward(x)
```

The forward function for the BaseModel.

Parameters

x (*Tensor*) – The inputs to the model.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a NotImplementedError.

```
class models.pytorch.models.YetAnotherTransformer(*args: Any, **kwargs: Any)
```

12 YetAnotherTransformer

A YetAnotherTransformer model that extends the Pytorch BaseModel.

Functionality: `#`. The class initializes with a set of parameters for the YetAnotherTransformerClassifier. `#`. It sets up the YetAnotherTransformerClassifier model, the optimizer and the learning rate scheduler. `#`. It also defines a forward method that applies the YetAnotherTransformerClassifier model.

12.1 Attributes

learning_rate

[float] The learning rate for the optimizer.

model

[YetAnotherTransformerClassifier] The YetAnotherTransformerClassifier model.

optimizer

[torch.optim.AdamW] The AdamW optimizer.

scheduler

[torch.optim.lr_scheduler.ExponentialLR] The learning rate scheduler.

12.2 Methods

forward(x)

Performs a forward pass through the model.

Args:

kwargs (dict): A dictionary containing the parameters for the YetAnotherTransformerClassifier, optimizer and learning rate scheduler.

Returns: None

```
__init__(**kwargs)
```

```
forward(x)
```

The forward function for the BaseModel.

Parameters

x (*Tensor*) – The inputs to the model.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a NotImplementedError.

```
class models.pytorch.models.YetAnotherTransformerClassifier(*args: Any, **kwargs: Any)
```

13 YetAnotherTransformerClassifier

A YetAnotherTransformerClassifier module that extends the nn.Module from PyTorch.

Functionality: `#`. The class initializes with a set of parameters for the transformer blocks. `#`. It sets up the transformer blocks and the output layer. `#`. It also defines a forward method that applies the transformer blocks, takes the mean over the time dimension of the transformed sequence, and applies the output layer.

13.1 Attributes

DEFAULTS

[dict] The default settings for the transformer.

settings

[dict] The settings for the transformer, with any user-provided values overriding the defaults.

transformer

[nn.ModuleList] The list of transformer blocks.

output_layer

[nn.Linear] The output layer.

13.2 Methods

forward(inputs)

Performs a forward pass through the model.

Args:

kwargs (dict): A dictionary containing the parameters for the transformer blocks.

Returns: None

`__init__(**kwargs)`

forward(inputs)

Forward pass through the model

Appendix L

Tensorflow Models

This module defines a PyTorch *BaseModel* providing a basic framework for learning and validating from *Trainer*

```
class models.tensorflow.models.BaseModel(*args: Any, **kwargs: Any)
```

A BaseModel that extends the tf.keras.Model.

Functionality: #. The class initializes with a given learning rate. #. It sets up the loss criterion, accuracy metric, and default states for optimizer and scheduler. #. It defines an abstract method 'call' which should be implemented in the subclass. #. It also defines various utility functions like calculating accuracy, training, validation and testing steps, scheduler stepping, and model checkpointing.

Args:

learning_rate (float): The initial learning rate for optimizer.

Parameters

learning_rate (*float*) – The initial learning rate for optimizer.

Returns

None

Return type

None

Note: The class does not directly initialize the optimizer and scheduler. They should be initialized in the subclass if needed.

Warning: The 'call' function must be implemented in the subclass, else it will raise a NotImplemented-Error.

```
__init__(learning_rate)
```

```
calculate_accuracy(y_pred, y_true)
```

Calculates the accuracy of the model's prediction.

Parameters

- y_pred (*Tensor*) – The predicted output from the model.
- y_true (*Tensor*) – The ground truth or actual labels.

Returns

The calculated accuracy.

Return type

float

```
call(inputs, training=False)
```

The call function for the BaseModel.

Parameters

- **inputs** (*Tensor*) – The inputs to the model.
- **training** (*bool*) – A flag indicating whether the model is in training mode. Default is False.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a `NotImplementedError`.**eval_mode()**

Sets the model to evaluation mode.

get_lr()

Gets the current learning rate of the model.

Returns

The current learning rate.

Return type

float

load_checkpoint(filepath)

Loads the model weights from a checkpoint.

Parameters**filepath** (*str*) – The file path where to load the model checkpoint from.**optimize()**

Sets the model to training mode.

save_checkpoint(filepath)

Saves the model weights to a checkpoint.

Parameters**filepath** (*str*) – The file path where to save the model checkpoint.**step_scheduler()**

Adjusts the learning rate according to the learning rate scheduler.

test_step(batch)

Performs a test step using the input batch data.

Parameters**batch** (*tuple*) – A tuple containing input data and labels.**Returns**

The calculated loss, accuracy, and model predictions.

Return type

tuple

train_mode()

Sets the model to training mode.

training_step(batch)

Performs a training step using the input batch data.

Parameters**batch** (*tuple*) – A tuple containing input data and labels.**Returns**

The calculated loss and accuracy.

Return type

tuple

`validation_step(batch)`

Performs a validation step using the input batch data.

Parameters

`batch` (*tuple*) – A tuple containing input data and labels.

Returns

The calculated loss and accuracy.

Return type

tuple

`class models.tensorflow.models.CVTransferLearningModel(*args: Any, **kwargs: Any)`

`__init__(**kwargs)`

`call(inputs, training=True)`

The call function for the BaseModel.

Parameters

- `inputs` (*Tensor*) – The inputs to the model.
- `training` (*bool*) – A flag indicating whether the model is in training mode. Default is False.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a `NotImplementedError`.

`class models.tensorflow.models.HybridModel(*args: Any, **kwargs: Any)`

`__init__(**kwargs)`

`call(inputs, training=True)`

The call function for the BaseModel.

Parameters

- `inputs` (*Tensor*) – The inputs to the model.
- `training` (*bool*) – A flag indicating whether the model is in training mode. Default is False.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a `NotImplementedError`.

`class models.tensorflow.models.LSTMClassifier(*args: Any, **kwargs: Any)`

LSTM-based Sequence Classifier

`__init__(**kwargs)`

`call(inputs)`

Forward pass through the model

`class models.tensorflow.models.LSTMPredictor(*args: Any, **kwargs: Any)`

`__init__(**kwargs)`

`call(inputs)`

The call function for the BaseModel.

Parameters

- `inputs` (*Tensor*) – The inputs to the model.

- **training** (*bool*) – A flag indicating whether the model is in training mode. Default is `False`.

Returns

`None`

Warning: This function must be implemented in the subclass, else it raises a `NotImplementedError`.

```
class models.tensorflow.models.TransformerEncoderLayer(*args: Any, **kwargs: Any)
```

A Transformer Encoder layer as a subclass of `tf.keras.layers.Layer`.

Functionality: `#`. The class first initializes with key parameters for `MultiHeadAttention` and feedforward network. `#`. Then it defines the key components like multi-head attention, feedforward network, layer normalization, and dropout. `#`. In the call function, it takes input and performs self-attention, followed by layer normalization and feedforward operation.

Args:

`d_model` (*int*): The dimensionality of the input. `n_head` (*int*): The number of heads in the multi-head attention. `dim_feedforward` (*int*): The dimensionality of the feedforward network model. `dropout` (*float*): The dropout value.

Parameters

- `d_model` (*int*) – The dimensionality of the input.
- `n_head` (*int*) – The number of heads in the multi-head attention.
- `dim_feedforward` (*int*) – The dimensionality of the feedforward network model.
- `dropout` (*float*) – The dropout value.

Returns

`None`

Return type

`None`

Note: The implementation is based on the “Attention is All You Need” paper.

Warning: Ensure that the input dimension ‘`d_model`’ is divisible by the number of attention heads ‘`n_head`’.

```
__init__(d_model, n_head, dim_feedforward, dropout)
```

```
class models.tensorflow.models.TransformerPredictor(*args: Any, **kwargs: Any)
```

A Transformer Predictor model that extends the `BaseModel`.

Functionality: `#`. The class first initializes with the learning rate and other parameters. `#`. It then creates an instance of `TransformerSequenceClassifier`. `#`. It also sets up the learning rate scheduler and the optimizer. `#`. In the call function, it simply runs the `TransformerSequenceClassifier`.

Args:

`kwargs` (*dict*): A dictionary of arguments.

Parameters

`kwargs` (*dict*) – A dictionary of arguments.

Returns

`None`

Return type

`None`

Note: The learning rate is set up with an exponential decay schedule.

Warning: The learning rate and gamma for the decay schedule must be specified in the 'kwargs'.

```
__init__(**kwargs)
```

```
call(inputs, training=True)
```

The call function for the BaseModel.

Parameters

- **inputs** (*Tensor*) – The inputs to the model.
- **training** (*bool*) – A flag indicating whether the model is in training mode. Default is False.

Returns

None

Warning: This function must be implemented in the subclass, else it raises a `NotImplementedError`.

```
class models.tensorflow.models.TransformerSequenceClassifier(*args: Any, **kwargs: Any)
```

A Transformer Sequence Classifier as a subclass of `tf.keras.Model`.

Functionality: `#`. The class first initializes with default or provided settings. `#`. Then it defines the key components like the transformer encoder layers and output layer. `#`. In the call function, it takes input and passes it through each transformer layer followed by normalization and dense layer for final output.

Args:

kwargs (dict): Any additional arguments. If not provided, defaults will be used.

Parameters

kwargs (*dict*) – Any additional arguments.

Returns

None

Return type

None

Note: The implementation is based on the “Attention is All You Need” paper.

Warning: The inputs should have a shape of (batch_size, seq_length, height, width), otherwise, a `ValueError` will be raised.

```
__init__(**kwargs)
```


Appendix M

Indices and tables

- [genindex](#)
- [modindex](#)
- [search](#)

Bibliography

- [1] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 2017.

Python Module Index

a

`augmentations`, [5](#)

C

`callbacks`, [7](#)

`config`, [9](#)

d

`data.data_utils`, [15](#)

`data.dataset`, [25](#)

`dl_utils`, [29](#)

m

`models.pytorch.models`, [39](#)

`models.tensorflow.models`, [51](#)

p

`predict_on_camera`, [23](#)

t

`trainer`, [33](#)

V

`visualizations`, [37](#)

Index

[\spxentry__getitem__\(\)](#) [\spxextradata.dataset.ASL_DATASET](#) [\spxentry__init__\(\)](#) [\spxextramodels.tensorflow.models.Transform](#)
 method, [25](#) method, [54](#)
[\spxentry__init__\(\)](#) [\spxextradata.dataset.ASL_DATASET](#) [\spxentry__init__\(\)](#) [\spxextramodels.tensorflow.models.Transform](#)
 method, [25](#) method, [55](#)
[\spxentry__init__\(\)](#) [\spxextradl_utils.DatasetWithLen](#) [\spxentry__init__\(\)](#) [\spxextramodels.tensorflow.models.Transform](#)
 method, [29](#) method, [55](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.BaseModel](#) [\spxentry__init__\(\)](#) [\spxextratrainer.Trainer](#) method,
 method, [40](#) [33](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.CVTransferLearningModel](#) [\spxextradl_utils.DatasetWithLen](#)
 method, [42](#) method, [29](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.HybridModel](#) [\spxextradata.dataset.ASL_DATASET](#)
 method, [43](#) method, [26](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.HybridModel](#) [\spxextradl_utils.DatasetWithLen](#)
 method, [43](#) method, [29](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.LSTMClassifier](#) [repr__\(\)](#) [\spxextradata.dataset.ASL_DATASET](#)
 method, [44](#) method, [26](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.LSTMPredictor](#) [weakref__](#) [\spxextradata.dataset.ASL_DATASET](#)
 method, [45](#) attribute, [27](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.MultiHeadSelfAttention](#) [\spxextradl_utils.DatasetWithLen](#)
 method, [45](#) attribute, [29](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.TransformerBlock](#) [weakref__](#) [\spxextratrainer.Trainer](#) at-
 method, [46](#) tribute, [34](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.TransformerEnsemble](#) [\spxentryadd_callback\(\)](#) [\spxextratrainer.Trainer](#)
 method, [47](#) method, [34](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.TransformerPredictor](#) [\spxentryASL_DATASET](#) [\spxextraclass](#) in
 method, [47](#) [data_loader](#), [25](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.TransformerSequenceClassifier](#) [\spxentryaugmentation_increase_callback\(\)](#) [\spxextrain](#)
 method, [48](#) module callbacks, [7](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.YetAnotherModel](#) [\spxentryaugmentations](#)
 method, [48](#) [\spxentrymodule](#), [5](#)
[\spxentry__init__\(\)](#) [\spxextramodels.pytorch.models.YetAnotherTransformer](#) [\spxentryBaseModel](#) [\spxextraclass](#) in mod-
 method, [50](#) els.pytorch.models, [39](#)
[\spxentry__init__\(\)](#) [\spxextramodels.tensorflow.models.BaseModel](#) [\spxentryBaseModel](#) [\spxextraclass](#) in mod-
 method, [51](#) els.tensorflow.models, [51](#)
[\spxentry__init__\(\)](#) [\spxextramodels.tensorflow.models.CVTransferLearningModel](#) [\spxentryBATCH_SIZE](#) [\spxextrain module config](#), [9](#)
 method, [53](#)
[\spxentry__init__\(\)](#) [\spxextramodels.tensorflow.models.HybridModel](#) [\spxentrycalculate_accuracy\(\)](#) [\spxextramodels.pytorch.models.Base](#)
 method, [53](#) method, [40](#)
[\spxentry__init__\(\)](#) [\spxextramodels.tensorflow.models.LSTMClassifier](#) [\spxentrycalculate_accuracy\(\)](#) [\spxextramodels.tensorflow.models.E](#)
 method, [53](#) method, [51](#)
[\spxentry__init__\(\)](#) [\spxextramodels.tensorflow.models.LSTMPredictor](#) [\spxentrycalculate_avg_landmark_positions\(\)](#) [\spxextrain](#)
 method, [53](#) module data.data_utils, [15](#)

`\spxentrycalculate_landmark_length_stats()\spxextrain` `\spxentryDYNAMIC_DROP_OUT_REDUCTION_RATE\spxextrain`
`module data.data_utils, 15` `module config, 10`
`\spxentrycall()\spxextramodels.tensorflow.models.BaseModel`
`method, 51` `\spxentryEARLY_STOP_METRIC\spxextrain mod-`
`\spxentrycall()\spxextramodels.tensorflow.models.CVTransferLearningModel` `module config, 10`
`method, 53` `\spxentryEARLY_STOP_MODE\spxextrain module`
`\spxentrycall()\spxextramodels.tensorflow.models.HybridModel` `config, 10`
`method, 53` `\spxentryEARLY_STOP_PATIENCE\spxextrain`
`\spxentrycall()\spxextramodels.tensorflow.models.LSTMClassifier` `module config, 10`
`method, 53` `\spxentryEARLY_STOP_TOLERANCE\spxextrain`
`\spxentrycall()\spxextramodels.tensorflow.models.LSTMPredictor` `module config, 10`
`method, 53` `\spxentryEPOCHS\spxextrain module config, 10`
`\spxentrycall()\spxextramodels.tensorflow.models.TransformerPredictor` `\spxentryeval_mode()\spxextramodels.pytorch.models.BaseModel`
`method, 55` `method, 40`
`\spxentrycallbacks` `\spxentryeval_mode()\spxextramodels.tensorflow.models.BaseModel`
`\spxentrymodule, 7` `method, 52`
`\spxentryCHECKPOINT_DIR\spxextrain module con-` `\spxentryevaluate()\spxextratrain.Trainer method, 34`
`fig, 9`
`\spxentryCLEANED_FILE\spxextrain module config,` `\spxentryFACE_FEATURE_START\spxextrain mod-`
`9` `ule config, 10`
`\spxentryCOLUMNS_TO_USE\spxextrain module` `\spxentryFACE_FEATURES\spxextrain module con-`
`config, 9` `fig, 10`
`\spxentryconfig` `\spxentryFACE_INDICES\spxextrain module config,`
`\spxentrymodule, 9` `10`
`\spxentryconvert_mp_to_df()\spxextrain module pre-` `\spxentryFACE_LANDMARKS\spxextrain module`
`dict_on_camera, 23` `config, 10`
`\spxentrycreate_data_loaders()\spxextrain module` `\spxentryforward()\spxextramodels.pytorch.models.BaseModel`
`data.data_utils, 16` `method, 40`
`\spxentryCVTransferLearningModel\spxextraclass in` `\spxentryforward()\spxextramodels.pytorch.models.CVTransferLea-`
`models.pytorch.models, 41` `method, 42`
`\spxentryCVTransferLearningModel\spxextraclass in` `\spxentryforward()\spxextramodels.pytorch.models.HybridEnsembl-`
`models.tensorflow.models, 53` `method, 43`
`\spxentrydata.data_utils` `\spxentryforward()\spxextramodels.pytorch.models.HybridModel`
`\spxentrymodule, 15` `method, 43`
`\spxentrydata.dataset` `\spxentryforward()\spxextramodels.pytorch.models.LSTMClassifier`
`\spxentrymodule, 25` `method, 44`
`\spxentryDATA_DIR\spxextrain module config, 9` `\spxentryforward()\spxextramodels.pytorch.models.LSTMPredictor`
`\spxentryDatasetWithLen\spxextraclass in dl_utils, 29` `method, 45`
`\spxentryDEVICE\spxextrain module config, 9` `\spxentryforward()\spxextramodels.pytorch.models.TransformerEn-`
`\spxentryDL_FRAMEWORK\spxextrain module con-` `method, 47`
`fig, 9` `\spxentryforward()\spxextramodels.pytorch.models.TransformerPre-`
`\spxentrydl_utils` `method, 47`
`\spxentrymodule, 29` `\spxentryforward()\spxextramodels.pytorch.models.TransformerSec-`
`\spxentrydropout_callback()\spxextrain module call-` `method, 48`
`backs, 8` `\spxentryforward()\spxextramodels.pytorch.models.YetAnotherEns-`
`\spxentryDYNAMIC_AUG_INC_INTERVAL\spxextrain` `method, 49`
`module config, 9` `\spxentryforward()\spxextramodels.pytorch.models.YetAnotherTra-`
`\spxentryDYNAMIC_AUG_INC_RATE\spxextrain` `\spxentryforward()\spxextramodels.pytorch.models.YetAnotherTra-`
`module config, 9` `method, 50`
`\spxentryDYNAMIC_AUG_MAX_THRESHOLD\spxextrain` `\spxentryframe_dropout()\spxextrain module augmen-`
`module config, 9` `tations, 5`
`\spxentryDYNAMIC_DROP_OUT_INIT_RATE\spxextrain`
`module config, 9` `\spxentryget_dataloader()\spxextrain module dl_utils,`
`\spxentryDYNAMIC_DROP_OUT_MAX_THRESHOLD\spxextrain` `30,`
`module config, 9` `\spxentryget_dataset()\spxextrain module dl_utils, 30`
`\spxentryDYNAMIC_DROP_OUT_REDUCTION_INTERVAL\spxextrain` `\spxentryget_lr()\spxextramodels.pytorch.models.BaseModel`
`module config, 9` `method, 40`

<code>\spxentryget_lr()\spxextramodels.tensorflow.models.BaseModel</code> method, 52	<code>\spxentryMIN_SEQUENCES\spxextrain</code> module config, 11
<code>\spxentryget_model_params()\spxextrain</code> module <code>dl_utils</code> , 31	<code>\spxentrymirror_landmarks()\spxextrain</code> module augmentations, 5
<code>\spxentryget_PT_Dataset()\spxextrain</code> module <code>dl_utils</code> , 30	<code>\spxentryMODEL_DIR\spxextrain</code> module config, 11
<code>\spxentryget_TF_Dataset()\spxextrain</code> module <code>dl_utils</code> , 30	<code>\spxentryMODELNAME\spxextrain</code> module config, 11
<code>\spxentryHAND_FEATURES\spxextrain</code> module config, 10	<code>\spxentrymodels.pytorch.models</code> <code>\spxentrymodule</code> , 39
<code>\spxentryHAND_INDICES\spxextrain</code> module config, 10	<code>\spxentrymodels.tensorflow.models</code> <code>\spxentrymodule</code> , 51
<code>\spxentryHybridEnsembleModel\spxextraclass</code> in <code>models.pytorch.models</code> , 42	<code>\spxentrymodule</code> <code>\spxentryaugmentations</code> , 5
<code>\spxentryHybridModel\spxextraclass</code> in <code>models.pytorch.models</code> , 43	<code>\spxentrycallbacks</code> , 7
<code>\spxentryHybridModel\spxextraclass</code> in <code>models.tensorflow.models</code> , 53	<code>\spxentryconfig</code> , 9
<code>\spxentryINPUT_SIZE\spxextrain</code> module config, 10	<code>\spxentrydata.data_utils</code> , 15
<code>\spxentryINTEREMOLATE_MISSING\spxextrain</code> module config, 10	<code>\spxentrydata.dataset</code> , 25
<code>\spxentryinterpolate_missing_values()\spxextrain</code> module <code>data.data_utils</code> , 16	<code>\spxentrydl_utils</code> , 29
<code>\spxentryLANDMARK_FILES\spxextrain</code> module config, 10	<code>\spxentrymodels.pytorch.models</code> , 39
<code>\spxentryLEARNING_RATE\spxextrain</code> module config, 10	<code>\spxentrymodels.tensorflow.models</code> , 51
<code>\spxentryLEFT_HAND_FEATURE_START\spxextrain</code> module config, 10	<code>\spxentrypredict_on_camera</code> , 23
<code>\spxentryLEFT_HAND_INDICES\spxextrain</code> module config, 10	<code>\spxentrytrainer</code> , 33
<code>\spxentryload_checkpoint()\spxextramodels.pytorch.models.BaseModel</code> method, 40	<code>\spxentryvisualizations</code> , 37
<code>\spxentryload_checkpoint()\spxextramodels.tensorflow.models.BaseModel</code> method, 52	<code>\spxentryMultiHeadSelfAttention\spxextraclass</code> in <code>models.pytorch.models</code> , 45
<code>\spxentryload_checkpoint()\spxextramodels.tensorflow.models.BaseModel</code> method, 52	<code>\spxentryN_CLASSES\spxextrain</code> module config, 11
<code>\spxentryload_data()\spxextradata.dataset.ASL_DATASET</code> method, 27	<code>\spxentryN_DIMS\spxextrain</code> module config, 11
<code>\spxentryload_relevant_data_subset()\spxextrain</code> module <code>data.data_utils</code> , 17	<code>\spxentryN_LANDMARKS\spxextrain</code> module config, 11
<code>\spxentrylog_metrics()\spxextrain</code> module <code>dl_utils</code> , 31	<code>\spxentrynormalize()\spxextrain</code> module augmentations, 5
<code>\spxentryLSTMClassifier\spxextraclass</code> in <code>models.pytorch.models</code> , 44	<code>\spxentryoptimize()\spxextramodels.pytorch.models.BaseModel</code> method, 40
<code>\spxentryLSTMClassifier\spxextraclass</code> in <code>models.tensorflow.models</code> , 53	<code>\spxentryoptimize()\spxextramodels.tensorflow.models.BaseModel</code> method, 52
<code>\spxentryLSTMPredictor\spxextraclass</code> in <code>models.pytorch.models</code> , 44	<code>\spxentryOUT_DIR\spxextrain</code> module config, 11
<code>\spxentryLSTMPredictor\spxextraclass</code> in <code>models.tensorflow.models</code> , 53	<code>\spxentryPOSE_FEATURE_START\spxextrain</code> module config, 11
<code>\spxentryMAP_JSON_FILE\spxextrain</code> module config, 11	<code>\spxentryPOSE_FEATURES\spxextrain</code> module config, 11
<code>\spxentryMARKER_FILE\spxextrain</code> module config, 11	<code>\spxentryPOSE_INDICES\spxextrain</code> module config, 11
<code>\spxentryMAX_SEQUENCES\spxextrain</code> module config, 11	<code>\spxentrypredict_on_camera</code> <code>\spxentrymodule</code> , 23
	<code>\spxentrypreprocess_data()\spxextrain</code> module <code>data.data_utils</code> , 17
	<code>\spxentrypreprocess_data_item()\spxextrain</code> module <code>data.data_utils</code> , 17
	<code>\spxentrypreprocess_data_to_same_size()\spxextrain</code> module <code>data.data_utils</code> , 18
	<code>\spxentrypreprocess_raw_data()\spxextrain</code> module <code>data.data_utils</code> , 18
	<code>\spxentryPROCESSED_DATA_DIR\spxextrain</code> module config, 11
	<code>\spxentryrandom_rotation()\spxextrain</code> module augmentations, 5

[\spxentryrandom_scaling\(\)\spxextrain module augmentations, 5](#)
[\spxentryRAW_DATA_DIR\spxextrain module config, 11](#)
[\spxentryremove_outlier_or_missing_data\(\)\spxextrain module data.data_utils, 19](#)
[\spxentryremove_unusable_data\(\)\spxextrain module data.data_utils, 19](#)
[\spxentryRIGHT_HAND_FEATURE_START\spxextrain module config, 11](#)
[\spxentryRIGHT_HAND_INDICES\spxextrain module config, 11](#)
[\spxentryROOT_PATH\spxextrain module config, 11](#)
[\spxentryROWS_PER_FRAME\spxextrain module config, 11](#)
[\spxentryRUNS_DIR\spxextrain module config, 11](#)
[\spxentrysave_checkpoint\(\)\spxextramodels.pytorch.models.BaseModel method, 41](#)
[\spxentrysave_checkpoint\(\)\spxextramodels.tensorflow.models.BaseModel method, 52](#)
[\spxentrySEED\spxextrain module config, 12](#)
[\spxentryshift_landmarks\(\)\spxextrain module augmentations, 5](#)
[\spxentryshow_camera_feed\(\)\spxextrain module predict_on_camera, 23](#)
[\spxentrySKIP_CONSECUTIVE_ZEROS\spxextrain module config, 12](#)
[\spxentrySRC_DIR\spxextrain module config, 12](#)
[\spxentrystandardize\(\)\spxextrain module augmentations, 6](#)
[\spxentrystep_scheduler\(\)\spxextramodels.pytorch.models.BaseModel method, 41](#)
[\spxentrystep_scheduler\(\)\spxextramodels.tensorflow.models.BaseModel method, 52](#)
[\spxentrytest\(\)\spxextratrainer.Trainer method, 35](#)
[\spxentryTEST_SIZE\spxextrain module config, 12](#)
[\spxentrytest_step\(\)\spxextramodels.pytorch.models.BaseModel method, 41](#)
[\spxentrytest_step\(\)\spxextramodels.tensorflow.models.BaseModel method, 52](#)
[\spxentryto_PT_DataLoader\(\)\spxextrain module dl_utils, 31](#)
[\spxentryto_TF_DataLoader\(\)\spxextrain module dl_utils, 31](#)
[\spxentrytrain\(\)\spxextratrainer.Trainer method, 35](#)
[\spxentryTRAIN_CSV_FILE\spxextrain module config, 12](#)
[\spxentrytrain_mode\(\)\spxextramodels.pytorch.models.BaseModel method, 41](#)
[\spxentrytrain_mode\(\)\spxextramodels.tensorflow.models.BaseModel method, 52](#)
[\spxentryTRAIN_SIZE\spxextrain module config, 12](#)
[\spxentrytrainer
 \spxentrymodule, 33](#)
[\spxentryTrainer\spxextraclass in trainer, 33](#)
[\spxentrytraining_step\(\)\spxextramodels.pytorch.models.BaseModel method, 41](#)
[\spxentrytraining_step\(\)\spxextramodels.tensorflow.models.BaseModel method, 52](#)
[\spxentryTransformerBlock\spxextraclass in models.pytorch.models, 45](#)
[\spxentryTransformerEncoderLayer\spxextraclass in models.tensorflow.models, 54](#)
[\spxentryTransformerEnsemble\spxextraclass in models.pytorch.models, 46](#)
[\spxentryTransformerPredictor\spxextraclass in models.pytorch.models, 47](#)
[\spxentryTransformerPredictor\spxextraclass in models.tensorflow.models, 54](#)
[\spxentryTransformerSequenceClassifier\spxextraclass in models.pytorch.models, 48](#)
[\spxentryTransformerSequenceClassifier\spxextraclass in models.tensorflow.models, 55](#)
[\spxentryTUNE_HP\spxextrain module config, 12](#)
[\spxentryUSED_FACE_FEATURES\spxextrain module config, 12](#)
[\spxentryUSED_HAND_FEATURES\spxextrain module config, 12](#)
[\spxentryUSED_POSE_FEATURES\spxextrain module config, 12](#)
[\spxentryUSEFUL_ALL_LANDMARKS\spxextrain module config, 12](#)
[\spxentryUSEFUL_FACE_LANDMARKS\spxextrain module config, 12](#)
[\spxentryUSEFUL_HAND_LANDMARKS\spxextrain module config, 12](#)
[\spxentryUSEFUL_LEFT_HAND_LANDMARKS\spxextrain module config, 12](#)
[\spxentryUSEFUL_POSE_LANDMARKS\spxextrain module config, 12](#)
[\spxentryUSEFUL_RIGHT_HAND_LANDMARKS\spxextrain module config, 12](#)
[\spxentryVALID_SIZE\spxextrain module config, 12](#)
[\spxentryvalidation_step\(\)\spxextramodels.pytorch.models.BaseModel method, 41](#)
[\spxentryvalidation_step\(\)\spxextramodels.tensorflow.models.BaseModel method, 52](#)
[\spxentryvisualizations
 \spxentrymodule, 37](#)
[\spxentryvisualize_data_distribution\(\)\spxextrain module visualizations, 37](#)
[\spxentryvisualize_target_sign\(\)\spxextrain module visualizations, 37](#)
[\spxentryYetAnotherEnsemble\spxextraclass in models.pytorch.models, 48](#)
[\spxentryYetAnotherTransformer\spxextraclass in models.pytorch.models, 49](#)
[\spxentryYetAnotherTransformerClassifier\spxextraclass in models.pytorch.models, 49](#)